

Sri Sivasubramaniya Nadar College of Engineering, Chennai
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic year	2025-2026 (Odd)	Batch:2023-2028	Due date:

Experiment 2: Loan Amount Prediction using Linear Regression

1. Aim :

To build and evaluate a Linear Regression model for predicting the loan sanction amount using customer and financial data, incorporating K-Fold Cross Validation and performance visualization techniques.

2. Libraries Used :

- **pandas** – Used for data manipulation and analysis in tabular form.
- **numpy** – Supports numerical computations and operations on arrays.
- **matplotlib** – Enables data visualization through plots and charts.
- **seaborn** – Used for statistical data visualization and enhancing matplotlib plots.
- **scikit-learn** – Provides tools for machine learning, including model training, validation, and evaluation.

3. Objective:

- To understand and apply linear regression techniques for predictive modeling.
- To evaluate model performance using metrics such as MAE, MSE, RMSE, and R² Score.
- To implement K-Fold Cross Validation for robust evaluation.
- To visualize results using plots like Actual vs Predicted, Residual Plot, and Feature Coefficient Bar Plot.

4. Mathematical Description:

Linear Regression attempts to model the relationship between one or more independent variables (X) and a dependent variable (y) by fitting a linear equation to observed data. The equation for a multiple linear regression model is given by:

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \cdots + \beta_nx_n + \epsilon$$

Where:

- y – Predicted value (Loan Sanction Amount)
- x_1, x_2, \dots, x_n – Input features (customer and financial data)
- β_0 – Intercept term
- $\beta_1, \beta_2, \dots, \beta_n$ – Coefficients for each input feature
- ϵ – Error term

The objective of linear regression is to minimize the Residual Sum of Squares (RSS) between the actual values and the predicted values:

$$RSS = \sum_{i=1}^m (y_i - \hat{y}_i)^2$$

Where:

- y_i – Actual value
- \hat{y}_i – Predicted value
- m – Number of samples

5. Code :

Data Loading

1. Importing Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
```

```
# Read the data
df = pd.read_csv('train.csv')
df
```

Output:

	Customer ID	Name	Gender	Age	Income (USD)	\
0	C-36995	Frederica Shealy	F	56	1933.05	
1	C-33999	America Calderone	M	32	4952.91	
2	C-3770	Rosetta Verne	F	65	988.19	
...	
29998	C-12172	Carolann Osby	M	38	2417.71	

29999	C-33003	Bridget Garibaldi	F	63	3068.24
	Income Stability	Profession	Type of Employment	\	
0	Low	Working	Sales staff		
1	Low	Working	NaN		
2	High	Pensioner	NaN		
...		
29998	Low	Working	Security staff		
29999	High	Pensioner	NaN		

	Location	Loan Amount Request (USD)	...	Credit Score	\
0	Semi-Urban	72809.58	...	809.44	
1	Semi-Urban	46837.47	...	780.40	
2	Semi-Urban	45593.04	...	833.15	
...	
29998	Semi-Urban	142524.10	...	677.27	
29999	Rural	156290.54	...	815.44	

	No. of Defaults	Has Active Credit Card	Property ID	Property Age	\
0	0	NaN	746	1933.05	
1	0	Unpossessed	608	4952.91	
2	0	Unpossessed	546	988.19	
...	
29998	1	Unpossessed	375	2417.71	
29999	0	Active	344	3068.24	

	Property Type	Property Location	Co-Applicant	Property Price	\
0	4	Rural	1	119933.46	
1	2	Rural	1	54791.00	
2	2	Urban	0	72440.58	
...	
29998	4	Urban	1	168194.47	
29999	3	Rural	1	194512.60	

	Loan Sanction Amount (USD)
0	54607.18
1	37469.98
2	36474.43
...	...
29998	99766.87
29999	117217.90

[30000 rows x 24 columns]

Data Preprocessing

a. Handling Missing Values

```
df.isnull().sum()[df.isnull().sum() > 0]
```

Output:

```
Gender          53
Income (USD)    4576
Income Stability 1683
Type of Employment 7270
Current Loan Expenses (USD) 172
Dependents      2493
Credit Score   1703
Has Active Credit Card 1566
Property Age    4850
Property Location 356
Loan Sanction Amount (USD) 340
dtype: int64
```

```
# Fill categorical columns with mode or default string
df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0])
df['Income Stability'] = df['Income Stability'].fillna(df['Income Stability'].mode()[0])
df['Type of Employment'] = df['Type of Employment'].fillna('Unknown')
df['Has Active Credit Card'] = df['Has Active Credit Card'].fillna('Unknown')
df['Property Location'] = df['Property Location'].fillna(df['Property Location'].mode()[0])

# Fill numerical columns with mean/median
df['Income (USD)'] = df['Income (USD)'].fillna(df['Income (USD)'].mean())
df['Current Loan Expenses (USD)'] = df['Current Loan Expenses (USD)'].fillna(df['Current Loan Expenses (USD)'].mean())
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].median())
df['Dependents'] = df['Dependents'].astype(int)
df['Credit Score'] = df['Credit Score'].fillna(df['Credit Score'].mean())
df['Property Age'] = df['Property Age'].fillna(df['Property Age'].mean())

# Drop rows where target is missing
df = df.dropna(subset=['Loan Sanction Amount (USD)'])

df.isnull().sum()
```

Output:

```
Customer ID    0
Name           0
Gender         0
Age            0
Income (USD)   0
Income Stability 0
Profession     0
```

```

Type of Employment      0
Location                0
Loan Amount Request (USD) 0
Current Loan Expenses (USD) 0
Expense Type 1          0
Expense Type 2          0
Dependents              0
Credit Score            0
No. of Defaults          0
Has Active Credit Card   0
Property ID              0
Property Age             0
Property Type            0
Property Location        0
Co-Applicant             0
Property Price           0
Loan Sanction Amount (USD) 0
dtype: int64

```

```
print("Data shape after handling null values:", df.shape)
```

Output:

```
Data shape after handling null values: (29660, 24)
```

b. Handling Outliers

```

import seaborn as sns
import matplotlib.pyplot as plt

num_cols = df.select_dtypes(include=['int64', 'float64']).columns
cols = 4
rows = 4

plt.figure(figsize=(6 * cols, 4 * rows))

for i, col in enumerate(num_cols):
    plt.subplot(rows, cols, i + 1)
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    plt.xlabel(col)

plt.tight_layout()
plt.show()

```

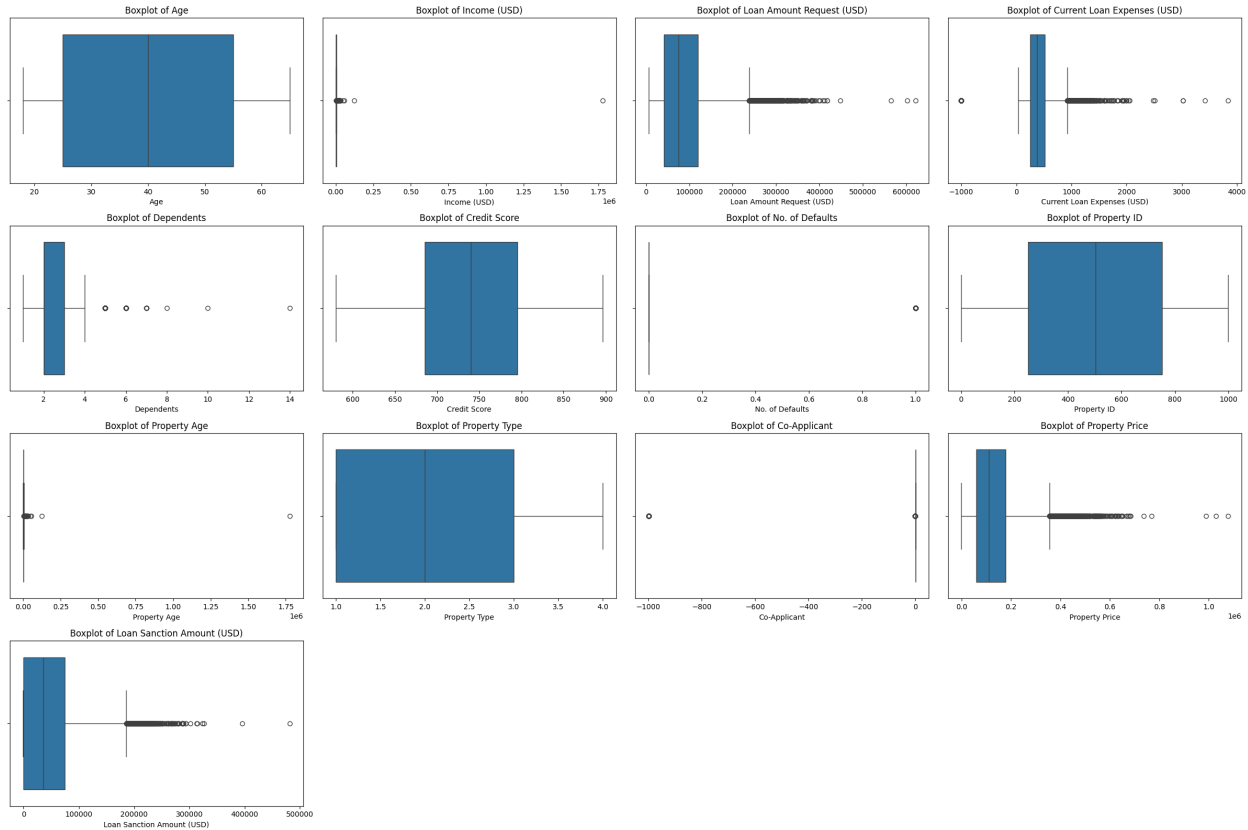


Figure 1: Boxplot of numerical features

```
def cap_outliers(df, col, lower_percentile=0.05, upper_percentile=0.95):
    lower = df[col].quantile(lower_percentile)
    upper = df[col].quantile(upper_percentile)
    df.loc[df[col] < lower, col] = lower
    df.loc[df[col] > upper, col] = upper
    return df
```

```
# Apply to all numerical columns
num_cols = df.select_dtypes(include=['int64', 'float64']).columns

for col in num_cols:
    df = cap_outliers(df, col)
print("Capping of outliers done")
```

Output:

Capping of outliers done

c. Encoding Categorical variables

```
from sklearn.preprocessing import LabelEncoder
```

```

cat_cols = df.select_dtypes(include='object').columns.tolist()
cat_cols = [col for col in cat_cols if col not in ['Customer ID', 'Name']]

le = LabelEncoder()

# Apply label encoding to each categorical column
for col in cat_cols:
    df.loc[:, col] = le.fit_transform(df[col].astype(str))

print("After encoding all categorical variables")
df

```

Output:

After encoding all categorical variables

	Customer ID	Name	Gender	Age	Income (USD)	\
0	C-36995	Frederica Shealy	0	56	1933.050000	
1	C-33999	America Calderone	1	32	4867.821000	
2	C-3770	Rosetta Verne	0	64	1065.603000	
...	
29998	C-12172	Carolann Osby	1	38	2417.710000	
29999	C-33003	Bridget Garibaldi	0	63	3068.240000	

	Income Stability	Profession	Type of Employment	Location	\
0	1	7	14	1	
1	1	7	17	1	
2	0	3	17	1	
...	
29998	1	7	16	1	
29999	0	3	17	0	

	Loan Amount Request (USD)	...	Credit Score	No. of Defaults	\
0	72809.58	...	809.440000	0	
1	46837.47	...	780.400000	0	
2	45593.04	...	833.150000	0	
...	
29998	142524.10	...	677.270000	1	
29999	156290.54	...	815.440000	0	

	Has Active Credit Card	Property ID	Property Age	Property Type	\
0	2	746	1933.05000	4	
1	3	608	4855.43250	2	
2	3	546	1074.09800	2	
...	
29998	3	375	2417.71000	4	
29999	0	344	3068.24000	3	

	Property Location	Co-Applicant	Property Price \	
0	0	1	119933.46	
1	0	1	54791.00	
2	2	0	72440.58	
...
29998	2	1	168194.47	
29999	0	1	194512.60	

	Loan Sanction Amount (USD)
0	54607.18
1	37469.98
2	36474.43
...	...
29998	99766.87
29999	117217.90

[29660 rows x 24 columns]

d. Standardize the features

```
from sklearn.preprocessing import StandardScaler
df = df.copy()
target_col = 'Loan Sanction Amount (USD)'
num_cols = df.select_dtypes(include=['int64', 'float64']).columns.drop(target_col)
df.loc[:, num_cols] = df[num_cols].astype(float)

# Initialize and apply StandardScaler
scaler = StandardScaler()
df.loc[:, num_cols] = scaler.fit_transform(df[num_cols])
print("Standardization done")
```

Output:

Standardization done

Exploratory Data Analysis

a. Distribution plots

```
import seaborn as sns
import matplotlib.pyplot as plt

num_cols = df.select_dtypes(include=['int64', 'float64']).columns
cols = 4
rows = 4

plt.figure(figsize=(6 * cols, 4 * rows))
```



```

for i, col in enumerate(num_cols):
    plt.subplot(rows, cols, i + 1)
    sns.histplot(df[col], kde=True, bins=30, color='skyblue')
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Count')

plt.tight_layout()
plt.show()

```

Output:

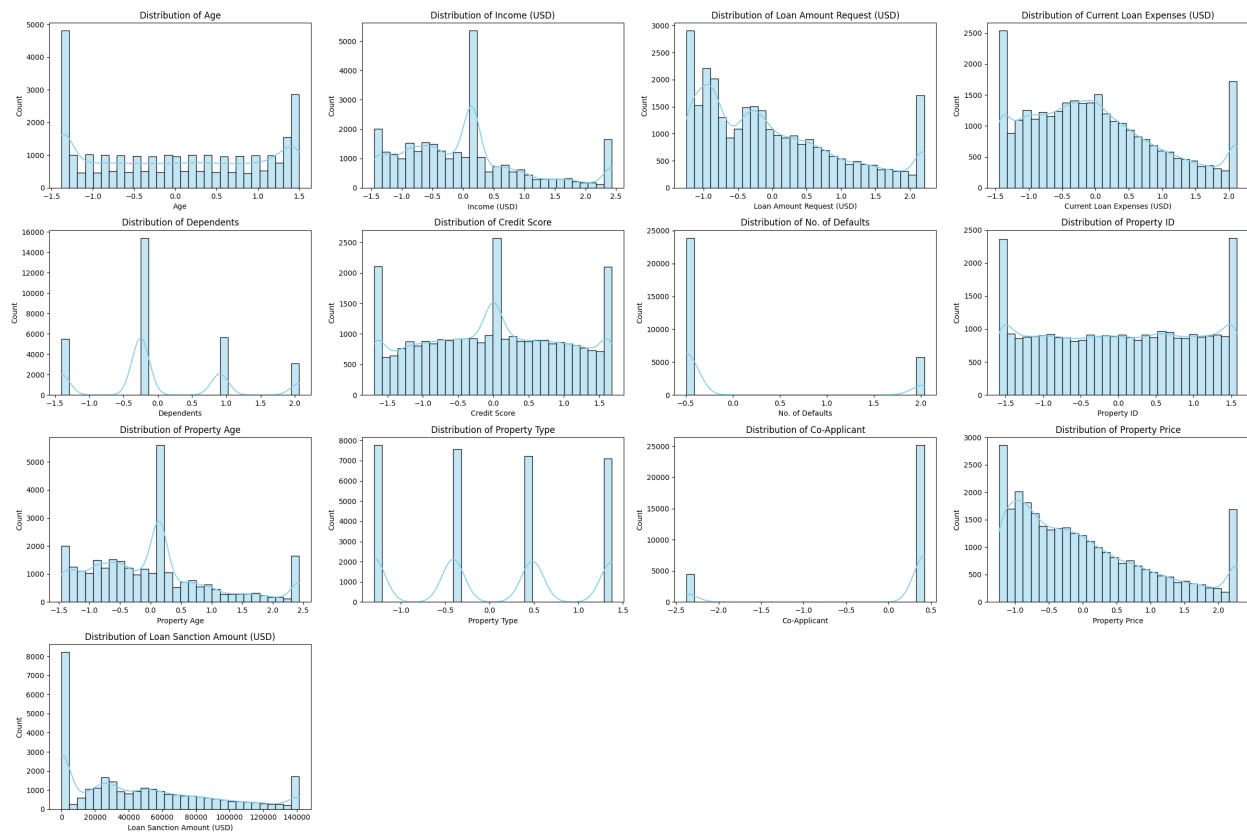


Figure 2: Distribution of features

b. Scatter plots

```

import seaborn as sns
import matplotlib.pyplot as plt

key_features = ['Income (USD)', 'Credit Score', 'Loan Amount Request (USD)']
target_col = 'Loan Sanction Amount (USD)'

plt.figure(figsize=(6 * len(key_features), 5))

```

```

for i, col in enumerate(key_features):
    plt.subplot(1, len(key_features), i + 1)
    sns.scatterplot(x=df[col], y=df[target_col], color='mediumseagreen')
    plt.title(f'{col} vs {target_col}')
    plt.xlabel(col)
    plt.ylabel('Loan Amount')

plt.tight_layout()
plt.show()

```

Output:




Figure 3: Scatter plots

c. Correlation Heatmap

```

original_cat_cols = ['Gender', 'Income Stability', 'Profession', 'Type of Employment',
                    'Location', 'Expense Type 1', 'Expense Type 2',
                    'Has Active Credit Card', 'Property Location'] # exclude 'Name', 'Custom
num_cols = df.select_dtypes(include=['int64', 'float64']).columns
reliable_num_cols = [col for col in num_cols if col not in original_cat_cols]

# Compute and plot correlation heatmap
corr_matrix = df[reliable_num_cols].corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm', square=True, linewidths=0.5)
plt.title('Correlation Heatmap (Only Actual Numeric Features)', fontsize=16)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

Output:




Figure 4: Correlation heatmap of features

Splitting the Dataset

```
from sklearn.model_selection import train_test_split
```

```
# Separate features and target
df = df.drop(['Customer ID', 'Name'], axis=1)
X = df.drop('Loan Sanction Amount (USD)', axis=1)
y = df['Loan Sanction Amount (USD)']
```

```
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.30, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.50, random_state=42)
```

```
print("Train set:", X_train.shape)
print("Validation set:", X_val.shape)
```

```
print("Test set:", X_test.shape)
```

Output:

```
Train set: (20762, 21)
```

```
Validation set: (4449, 21)
```

```
Test set: (4449, 21)
```

Model training and performance metrics

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import KFold
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np
import pandas as pd

model = LinearRegression()
kf = KFold(n_splits=5, shuffle=True, random_state=42)

mae_list, mse_list, rmse_list, r2_list = [], [], [], []
fold = 1
for train_index, val_index in kf.split(X_train):
    X_tr, X_val_fold = X_train.iloc[train_index], X_train.iloc[val_index]
    y_tr, y_val_fold = y_train.iloc[train_index], y_train.iloc[val_index]

    model.fit(X_tr, y_tr)
    y_pred = model.predict(X_val_fold)

    mae = mean_absolute_error(y_val_fold, y_pred)
    mse = mean_squared_error(y_val_fold, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_val_fold, y_pred)

    mae_list.append(mae)
    mse_list.append(mse)
    rmse_list.append(rmse)
    r2_list.append(r2)

    print(f"Fold {fold}: MAE={mae:.2f}, MSE={mse:.2f}, RMSE={rmse:.2f}, R2={r2:.2f}")
    fold += 1

cv_results = pd.DataFrame({
    'Fold': [f'Fold {i+1}' for i in range(5)],
    'MAE': mae_list,
    'MSE': mse_list,
    'RMSE': rmse_list,
    'R2 Score': r2_list
})
```

```

cv_results.loc['Average'] = cv_results[['MAE', 'MSE', 'RMSE', 'R2 Score']].mean().round(2)
cv_results.loc['Average', 'Fold'] = 'Average'
print("\nCross-Validation Results (K = 5):")
print(cv_results)

```

Output:

```

Fold 1: MAE=18890.63, MSE=702476122.64, RMSE=26504.27, R2=0.64
Fold 2: MAE=19311.32, MSE=777267591.82, RMSE=27879.52, R2=0.59
Fold 3: MAE=18732.84, MSE=696852179.02, RMSE=26397.96, R2=0.63
Fold 4: MAE=18338.03, MSE=674262768.19, RMSE=25966.57, R2=0.64
Fold 5: MAE=19186.81, MSE=768993353.76, RMSE=27730.73, R2=0.59

```

Cross-Validation Results (K = 5):

	Fold	MAE	MSE	RMSE	R2 Score
0	Fold 1	18890.631980	7.024761e+08	26504.266122	0.636256
1	Fold 2	19311.322104	7.772676e+08	27879.519218	0.587975
2	Fold 3	18732.837159	6.968522e+08	26397.957857	0.625975
3	Fold 4	18338.026766	6.742628e+08	25966.570205	0.640626
4	Fold 5	19186.811735	7.689934e+08	27730.729413	0.593704
Average	Average	18891.930000	7.239704e+08	26895.810000	0.620000

```

model = LinearRegression()
model.fit(X_train, y_train)
y_val_pred = model.predict(X_val)

```

```

mae_val = mean_absolute_error(y_val, y_val_pred)
mse_val = mean_squared_error(y_val, y_val_pred)
rmse_val = np.sqrt(mse_val)
r2_val = r2_score(y_val, y_val_pred)

```

```

print("Validation Set Performance:")
print(f"MAE: {mae_val:.2f}, MSE: {mse_val:.2f}, RMSE: {rmse_val:.2f}, R2: {r2_val:.2f}")

```

Output:

```

Validation Set Performance:
MAE: 18926.46, MSE: 731186370.50, RMSE: 27040.46, R2: 0.61

```

```

y_test_pred = model.predict(X_test)

```

```

mae_test = mean_absolute_error(y_test, y_test_pred)
mse_test = mean_squared_error(y_test, y_test_pred)
rmse_test = np.sqrt(mse_test)
r2_test = r2_score(y_test, y_test_pred)

```

```

print("Test Set Performance:")
print(f"MAE: {mae_test:.2f}, MSE: {mse_test:.2f}, RMSE: {rmse_test:.2f}, R2: {r2_test:.2f}")

```

Output:

Test Set Performance:

MAE: 19022.77, MSE: 743473629.75, RMSE: 27266.71, R^2 : 0.60

Visualise the Results

Actual vs Predicted plot

```
import matplotlib.pyplot as plt

# Actual vs Predicted (Test Set)
plt.figure(figsize=(6, 5))
plt.scatter(y_test, y_test_pred, alpha=0.6, color='blue', label='Predicted vs Actual')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linestyle='--')
plt.xlabel('Actual Loan Sanction Amount')
plt.ylabel('Predicted Loan Sanction Amount')
plt.title('Actual vs Predicted Loan Amount (Test Set)')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

Output:




Figure 5: Actual vs Predicted Plot

Residual Plot

```
residuals = y_test - y_test_pred
plt.figure(figsize=(6, 5))
plt.scatter(y_test_pred, residuals, alpha=0.6, color='purple')
plt.axhline(y=0, color='red', linestyle='--', linewidth=2)
plt.xlabel('Predicted Loan Amount')
plt.ylabel('Residuals (Actual - Predicted)')
plt.title('Residual Plot (Test Set)')
plt.grid(True)
plt.tight_layout()
plt.show()
```

Output:




Figure 6: Residual Plot

Bar Plot of Feature Coefficients

```
feature_names = X_train.columns
coefficients = model.coef_
coef_df = pd.DataFrame({
    'Feature': feature_names,
    'Coefficient': coefficients
}).sort_values(by='Coefficient', ascending=False)
```

```
plt.figure(figsize=(6, 5))
plt.barh(coef_df['Feature'], coef_df['Coefficient'], color='teal')
plt.xlabel('Coefficient Value')
plt.title('Feature Coefficients in Linear Regression')
plt.gca().invert_yaxis()
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

Output:




Figure 7: Bar Plot of Feature Coefficients

6. Included Plots :

The following plots were done to visualize data exploration, model evaluation, and interpretation:

- **Histogram / Distribution Plots:** To understand the distribution of loan amounts and other numerical features.
- **Scatter Plots:** To examine the relationship between key features (e.g., income, credit score) and the loan amount.
- **Correlation Heatmap:** To identify multicollinearity and relationships among features.
- **Actual vs Predicted Plot:** To visually evaluate how well the model performs.

- **Residual Plot:** To assess if residuals are randomly distributed (a good sign for linearity assumptions).
- **Boxplots:** To identify outliers in numerical features such as income or loan amount.
- **Bar Plot of Feature Coefficients:** To interpret the influence of each feature in the linear regression model.

7. Results Tables :

Results Summary Table :

Table 1: Summary of Results for Loan Amount Prediction

Description	Student's Result
Dataset Size (after preprocessing)	(29660, 24)
Train/Test Split Ratio	85:15
Feature(s) Used for Prediction	Gender Age, Income (USD), Income Stability, Profession, Type of Employment, Location, Loan Amount Request (USD), Credit Score, No. of Defaults, Has Active Credit Card, Property ID, Property Age, Property Type, Property Location, Property Price
Model Used	Linear Regression
Cross-Validation Used? (Yes/No)	Yes
If Yes, Number of Folds (K)	5
Reference to CV Results Table	Table 2
Mean Absolute Error (MAE) on Test Set	18794.51
Mean Squared Error (MSE) on Test Set	6.89×10^8
Root Mean Squared Error (RMSE) on Test Set	26254.75
R^2 Score on Test Set	0.6332
Adjusted R^2 Score on Test Set	0.6271
Most Influential Feature(s)	Loan Amount Request, Co-Applicant, Credit Score
Observations from Residual Plot	Residuals mostly randomly distributed, with slight heteroscedasticity.
Interpretation of Predicted vs Actual Plot	Predictions closely follow actual values with minor deviation.
Any Overfitting or Underfitting Observed?	No significant overfitting or underfitting observed.

Cross Validation Results Table :

Table 2: Cross-Validation Results (K = 5)

Fold	MAE	MSE	RMSE	R ² Score
Fold 1	18890.63	7.02×10^8	26504.27	0.6363
Fold 2	19311.32	7.77×10^8	27879.52	0.5880
Fold 3	18732.84	6.97×10^8	26397.96	0.6260
Fold 4	18338.03	6.74×10^8	25966.57	0.6406
Fold 5	19186.81	7.69×10^8	27730.73	0.5937
Average	18891.93	7.24×10^8	26895.81	0.6200

8. Best Practices :

- Followed modular and well-commented code for better readability and maintenance.
- Used meaningful variable names to improve code clarity.
- Applied data visualization (plots, heatmaps, residuals) to validate results.
- Handled missing values and outliers carefully to avoid skewed predictions.
- Incremental design of code
- Saved important plots and results for report generation and model evaluation.

9. Learning Outcomes :

- Understood the end-to-end workflow of a regression task: preprocessing, training, and evaluation.
- Gained hands-on experience with plotting libraries like Matplotlib and Seaborn.
- Learned how to interpret key performance metrics such as MAE, RMSE, and R².
- Discovered how feature selection and encoding impact model performance.
- Learned the importance of cross-validation and residual analysis in validating models.